Inside HPT

B5 / PullCube · 128 tokens × 256 dimensions

Heterogeneous Pretrained Transformer: modality-specific stems → shared transformer → task readouts

Observation-dependent features

RGB stem
16 × 256

State stem
16 × 256

History tokenizer + resampler
16 × 256 · B5 extension

Learned input tokens · EgoWAM

64 action tokens
Learned parameter vectors

16 future tokens
Learned parameter vectors

Concat.
+ position

Shared transformer trunk

16 blocks · width 256 · 8 attention heads

LayerNorm → self-attention
Residual addition

LayerNorm → MLP
Residual addition

All token groups attend to each other.

One forward pass per observation.

No noisy trajectories or flow time τ here.

Read out the corresponding positions

64 action-context tokens
Condition the action flow head

16 future-context tokens
Condition the trajectory decoder

Both outputs depend on this observation.

Token counts are latent capacity;
they do not count predicted time steps.

Two separate inputs to the trajectory decoder

HPT context is fixed during 50 generation steps

16 future-context tokens
From HPT

Condition via cross-attention

Trajectory decoder
Predict a velocity field

Update
the trajectory

Flow time τ

Noisy displacements + anchors
+ observed / noisy slot flags

Tokenizer Tθ
One token per point

The 16 HPT context tokens bypass Tθ. The decoder combines context with the evolving trajectory.